

ATL105 | October 7 Run4 Independent Review

REVIEW REQUIRED. Verified structural evidence is not semantic approval or host certification.

Measure	Recount
AI BRs	5153
AI TSs	14297
AI TC candidates	21531
AI BRs directly linked to candidate TCs (internal graph)	3719
TSs with TC	12860
TSs without TC	1437
Unique physical request filenames	21621
Physical request copies across candidate/reporting locations	23271
Logical TCs with verified physical metadata	21507
Missing declared TD files	24
Empty request bodies (physical copies)	764

Manual corpus search findings: 8; Run4-specific AI-team feedback is queued, not sent. All formal reviews remain PENDING_SEMANTIC_REVIEW, with no authorized SME decision or coverage credit.

Manual-search findings JSON

Independent Test Solution BR coverage: 0/601 confirmed Run4 matches (0.00% confirmed evidence), status NOT_ASSESSED; all rules remain unassessed. A complete source-anchor candidate crosswalk covers all 5,153 AI BRs: 562 AI BRs have candidate links to 270 distinct Test BRs, but candidates are not semantic confirmations. The matrix has 65,989 leaves; direct BR>TS>TC>physical-TD paths reach 3,719 BRs. Zero confirmed matches do not prove no equivalent AI BRs exist. The former 3,719/5,153 (72.17%) is AI-internal linkage only; full semantic coverage remains NOT_CALCULABLE.

Source-backed review workflow: 1,950 candidate pairs across 270 independent Test BRs; 331 require manual corpus search. All 1,950 Run4 pair decisions remain PENDING. No semantic confirmations or coverage credit have been granted.

Test Team interpretation for Element 4: both BRs are correct and semantically aligned. Element 4 is Address Line 2, character type AN, fixed width 21 bytes, carrying city/state/ZIP content; the Test BR's positional detail is compatible specificity, not a contradiction. The formal pair decision remains PENDING until an authorized reviewer records it; no coverage credit is issued.

Dated Test Team interpretation addendum

Section 10.1 independent Test Solution draft supplement: 59 BR -> TS -> TC -> TD design chains and 14 existing rule references. All draft chains remain REVIEW_REQUIRED; TD designs are non-executable; baseline unchanged; no coverage credit.

Section 10.1 chain assessment

Section 10.2 debit draft supplement: 29 debit-specific BR -> TS -> TC -> TD chains; 42 unique Section 10.1 references; 3 existing baseline references. All 27 source headings are structurally mapped; every reference remains review-required, with no coverage credit or execution certification.

Section 10.2 heading-to-chain assessment with IDs

Multi-leg ZIP: 406 files / 89 cases; CRC failure None; 0 missing and 0 differing candidate members. Package integrity is not flow or host certification.

HIGH | 764 empty request bodies counted across physical copies

This is a candidate/reporting copy count, not 764 distinct TCs. A present file is not a populated executable request; TC-000001 is an Auth Completion (0220) empty object.

Required action: Resolve message-family/template grounding or classify blocked; do not call empty bodies real executable data.

HIGH | 24 declared request-data files are missing

Catalog declares physical files not found; complete affected-case list is exported.

Required action: Producer supplies corrected payload/metadata pairs with declared case links.

HIGH | Existing independent Java adapter cannot map current payload roots

21,621 metadata records were read; element observations and field mappings are empty. All records remain REVIEW_REQUIRED.

Required action: Review a run-versioned root-family crosswalk and exact field identities before rerunning value rules; no implicit alias promotion.

HIGH | 1,650 producer approvals are automated-batch-nonSME

Approval audit count/reviewer identity verified; source report also cautions that approval is demo/sample rather than real SME approval.

Required action: Do not claim SME-approved or execution-ready suite; retain reviewer provenance and review required status.

HIGH | Reused IDs contain materially changed content

5,141 of 5,153 reused BR statements changed; all 21,123 reused TC fingerprints changed compared with October 5.

Required action: Do not migrate older matches by ID; remap by version/source/context and business meaning.

MEDIUM | TC audit reports zero cases

The actual catalog contains 21,531 TCs; the TC audit summary is not a reliable coverage gate.

Required action: Regenerate the audit from the final case catalog and pin hashes/timestamps.

MEDIUM | ASSERT figures use conflicting or unexplained counting units

Run report says 516; TD audit says 3,476; 8,637 cases contain ASSERT-tier nodes, with 37,497 ASSERT nodes and 6,710 cases carrying explicit assertion/value evidence under the analyzer's disclosed criterion.

Required action: Separate node counts, case counts and grounded/qualified assertions; do not sum overlapping tiers or treat ASSERT labels as certified oracles.

MEDIUM | Matrix proposes 62 BR edges absent from direct TC links

Full JSON has 65,989 leaves; 3,720 matrix-linked BRs versus 3,719 direct-linked BRs. Matrix-only BR is :748; relinks remain producer candidates.

Required action: Review each source-rule relink against underlying BR/scenario/context; preserve direct graph and proposed graph separately.

MEDIUM | 28 FULLY_TRACED leaves have no TD declaration

Status may be propagated from a BR with another complete chain; individual leaves are not proven complete.

Required action: Emit leaf-level and BR-level trace states separately; no automatic semantic or readiness pass.

MEDIUM | 1,437 scenarios have no TC

Generation backlog improves from 3,732 but remains in the denominator; no approved exclusion is established.

Required action: Use exported backlog for grounded regeneration or explicit reviewed scope decisions.

MEDIUM | 193 cases have no direct BR attribution

Matrix reports 131 orphan scenarios and documents 62 source-rule relinks; this is a proposed attribution distinction, not confirmed semantic matching.

Required action: Validate source/context justification for proposed edges and disclose residual unattributed artifacts.

HIGH | Business equivalence and host readiness remain unqualified

Expected-response structures now exist but contain INFER/GROUNDED_SECONDARY/FLAG claims; Judge bridge caveat, partial templates and root-mapping limitations remain.

Required action: Independently qualify rule interpretations, assertions, applicable data/wire and authorized host/state/response evidence; record SME questions without granting approval.